

AGENTIC-ARBITER


Free-cooling decision report · Union Pacific Jack Koraleski Technology, Research & Development Center, NE · station KBTA · 2023-03-30

An agent decides, hour by hour, whether this data centre can switch its mechanical chillers off and cool with outside air instead. Every hour it releases carries a safety margin measured from the agent's own past errors, and it declines the hours it cannot stand behind.

\$11k – \$47k at this site: \$11,393 to \$46,714, at \$3,216 to \$6,595 per MW of IT load per year MODELED	7.2% less mechanical cooling runtime, a share so it holds at any hall size	+236 h chiller-hours recovered per year against the control operators run today	43,230 hours of real recorded weather, 912 days held out	65.6% bound coverage on the four measured forecast pairs, against a 90% target
--	--	---	--	--

The agent's day, hour by hour

3 of 24 hours free cooling



What this day shows

- Just **3 of the 24 hours** ran on outside air. The remaining **21** ran chillers, and every one of them has a named reason rather than a default.
- The largest single block was **not weather**. The switch budget accounts for 14 of them, 00:00 to 13:00. The air was cold enough; the plant had already committed its allowance of mode changes for the day.
- The **bound held** across 912 days the agent never trained on. It took **14,469** free-cooling hours and **9** of them turned out unsafe: 0.62 per thousand.
- The **forecast is the product**: hold everything else and drop the forecast to no skill at all, and the 236 hours a year this recovers becomes 27.


What makes this different from a thermostat with a forecast attached: where the site geometry defeats the physics, the agent refuses the hour rather than return a number it cannot stand behind, and the cost of that caution is measured and published rather than hidden.

Contents

This document is one site's complete analysis. Page 1 is the whole argument; everything after it is the evidence.

1	Executive summary	what the agent did on this day, what it is worth, and how well the bound held
2	The site	the one mapped building at this site, and what the imagery can and cannot show
3	The findings	the bound against the limit hour by hour, and the comparison with the incumbent
4	How the decision is made	the margin and where each degree of it comes from, the gates, the switch budget
5	Validation, and how the bound is measured	held-out testing, and how the bound's coverage is measured as the record grows
6	Scale, and what it is worth	the same agent across every covered site, and who buys this
7	Appendix: every hour, and each reason once	the full schedule, and each distinct reason explained once

The site



One building, Union Pacific Jack Koraleski Technology, Research & Development Center. It is the only data centre footprint mapped at this location, so there is no neighbouring hall discharging heat across a shared yard and nothing rejecting warm air toward the intake this agent is protecting.

That matters for what this report contains. Where two halls face each other, the exhaust from one can reach the air the other breathes, and the agent has to solve that **plume** on the real footprints before it will release an hour. A single site has no such term: the plume physics does not apply here, so the section on it is absent from this document rather than present and empty.

Everything else is unchanged. The bound is still built from this site's own measured forecast error, still checked against the plant's own limit, and still declines any hour it cannot stand behind.


ESRI World Imagery. 0.3–0.5 m resolution shows objects, not nameplates. Cannot certify unit type or measure heights. The outline is OpenStreetMap way 166431620 on an independently georeferenced basemap, so an edge can sit a few metres off the roofline; every distance in this report is computed from the OpenStreetMap geometry, not from the picture.

The findings

Charts first. The prose under each one says only what the chart cannot.

The bound against the plant limit

degrees Celsius at the intake



The pale band is the safety margin: the distance the agent adds on top of the raw forecast before it will commit to an hour. The blue line is the resulting upper bound, the navy line is what the intake temperature actually did, and the dashed line is the plant's limit.

Read the crossings: the actual went above the bound in 7 of the 24 hours, which is what page 6 measures directly. All 7 were hours the agent had already put on chillers, and **0 of the hours it released for free cooling exceeded the plant limit**. A bound exceeded on an hour the agent did not act on costs nothing; the hours that matter are the ones it said yes to.

Chiller runtime, agent against incumbent

hours of mechanical cooling across 912 days the agent never trained on



7.2% less chiller runtime, a share so it holds at any hall size


The comparison is against a tuned reactive controller using an on-site sensor, which is what operators verifiably run today, not against doing nothing. **Held-out** means days the agent never trained on: the record was split chronologically and the second half, 912 days, was kept back.

How the decision is made

Each hour the agent takes the forecast for this building's own tile, adds a margin, and compares the result against the plant's limit. It commits to free cooling only when the **upper bound** clears the **limit**, not when the forecast does.

Where each degree of the margin comes from

degrees Celsius added on top of the forecast



Every bar here is forecast error alone. The plume allowance that would sit on top of it is identically zero at this site, because there is no neighbouring hall for an exhaust to cross from, so the chart has one part rather than two.

The margin is measured, not chosen. The lower part of each bar is the forecast error this agent has actually made at that hour of the day (**group-conditional**: the error is measured per hour of day rather than pooled across all of them, because a 3 a.m. forecast and a 3 p.m. forecast are not equally hard).

The gates an hour has to pass

Gate	Setting here	What it protects
Plant limit	18.0 °C	the intake temperature the hall is committed to
Dew point	15.0 °C	condensation on cold surfaces inside the hall
Notice	3 h	the plant cannot change mode instantly
Switch budget	1 per day	chillers and dampers wear out when cycled
Minimum dwell	3 h	a mode has to hold before another change is allowed


The switch budget is the reason a cold hour can still run chillers. With 1 change permitted in a day, spending one early costs the agent the ability to switch later, so it schedules the whole day at once by dynamic programming rather than deciding each hour as it arrives.

What this document is · the interface recomputes for any configuration a reader selects. This file was generated for the one named above, chosen as the configuration where this site's agent does the most while still changing mode at least once. If the screen disagrees with this page, the configuration differs.

Validation, and how the bound is measured

Bound coverage, by population size

share of hours the real intake stayed under the bound



A **bound** is a number the agent commits to the real intake staying under. Coverage is the share of hours it did. The target is 90%, and each population below is as large as the record currently allows.


What the measurement shows

- The method reaches its target **wherever the record is large enough to measure it**. Twelve separately calibrated per-lead bounds, every one at or above 90%, the worst at 91%. An adaptive bound over 42,727 rounds of the five-year record realised 90%.
- **The FortyGuard-specific calibration currently rests on 4 measured forecast days**, and reads 66%. A one-sided 90% bound needs **9** such days: at 4 the arithmetic ceiling is 80%, so the target is a function of how many days have been forecast rather than of the method.
- A calibration day is **earned, not bought**. It is a forecast plus the day that follows it, so the count rises by one for every further day the site is forecast on. Nine days of operation is the whole distance between the figure above and the target.
- Conditioning already recovers **most of the gap**. Measuring the error per hour of day rather than pooling every hour together lifts the worst group from 77% to 88%.

What was tested	Volume	Result
Real recorded weather	43,230 h / 1,824 d	the record the schedule was scored on
Held-out days	912 d	never trained on
Re-plans compared	21,684	94% changed nothing at all
Free-cooling hours taken	14,469	9 unsafe, 0.62 per thousand

Scale, and what it is worth

One hall is a pilot. The agent covers **238 sites**, each built on its own OpenStreetMap footprint, its own station weather record and its own tariff, drawn from 639 tagged data centre facilities mapped across the United States.



This is every site the agent is offered on. **12 more were measured and excluded** because the agent's own constraints made it worse than the incumbent there, and those are not sold without site-specific engineering work on them first: at that geometry the safety margin hands back more free-cooling hours than it wins, so the agent runs the chillers more than the controller it would replace, the worst of them by 3,649 hours a year. They are not in the chart, because an axis wide enough to hold them would compress everything above into a sixth of the frame; they are counted here instead, and they stay on the map marked measured rather than ready. A portfolio that reports where it does not work is one a reader can believe about where it does.

Who buys this

An operator running many data centres, and the engineer inside it who owns the energy number. Their tenants' service level agreements commit them to a maximum supply air temperature, so an excursion is a contract breach and not just a warm hour, which is why the conservative buffer exists in the first place. What replaces the buffer has to come with a bound and a published breach rate, or it does not get switched on.

At this site	Value	Basis
Measured footprint	5,049 m²	OpenStreetMap rings, the same outlines the solver runs on
IT load	4 to 7 MW	DERIVED from a watts per square metre density
Recovered hours	+236 h/yr	backtested
Value of those hours	\$11,393 to \$46,714 /yr	MODELED. Measured hours times published rates times derived megawatts
The rate behind it	\$3,216 to \$6,595 /MW-IT/yr	32 cells: published tariffs x published chiller efficiencies, swept

The percentage is the honest headline and the dollars are the illustration. 7.2% less chiller runtime needs no assumption about how big the building is; the money rows do, and the two that depend on one are labelled.

Appendix: every hour, and each reason once

The 24 hours of this day carry **3 distinct reasons** between them. Each is explained once below, with the hours it covers and the tightest of those hours as the worked example.

Blocked by the switch budget 14 hours: 00:00 to 13:00

These hours were **cold enough to free-cool and were run on chillers anyway**, and the reason is the schedule rather than the weather. The plant permits 1 mode change a day. Spending one here would leave the agent unable to switch when a longer or colder block arrives, so the whole day is planned at once and these hours are the ones the plan gives up. This is the explanation a thermostat cannot give, because a thermostat has no plan to be constrained by.

Worked example, 13:00 · ambient air **19.89 °C**, less the **3.44 °C** the forecast earns over assuming today repeats, plus a measured **1.39 °C** margin, gives a **17.84 °C** bound against the 18.0 °C limit. With no neighbouring hall to warm the incoming air, the intake is the ambient air, and it reached 19.89 °C.

Blocked by air temperature 7 hours: 14:00 to 20:00

The air was simply too warm. In each of these hours the agent's upper bound on intake temperature sat above the plant's 18.0 °C limit, so free cooling would have risked the hall. It is the bound that is compared against the limit, never the raw forecast: the margin is what makes the comparison safe rather than optimistic.

Worked example, 20:00 · ambient air **18.11 °C**, less the **0.16 °C** the forecast earns over assuming today repeats, plus a measured **1.35 °C** margin, gives a **19.3 °C** bound against the 18.0 °C limit. With no neighbouring hall to warm the incoming air, the intake is the ambient air, and it reached 18.11 °C.

Free cooling ran 3 hours: 21:00 to 23:00

Free cooling ran. In each of these hours the agent's upper bound cleared the plant limit with the full measured margin already added, so outside air did the cooling and the chillers stayed off.

Worked example, 21:00 · ambient air **15 °C**, plus **1.38 °C** where the forecast read warmer than assuming today repeats, plus a measured **1.32 °C** margin, gives a **17.7 °C** bound against the 18.0 °C limit. With no neighbouring hall to warm the incoming air, the intake is the ambient air, and it reached 15 °C.

The full schedule

hh	mode	safe	reason	bound	limit	actual	margin
00	mechanical	yes	switch budget	-0.61	18.0	-0.61	0.848
01	mechanical	yes	switch budget	-0.75	18.0	-1.0	0.822
02	mechanical	yes	switch budget	-0.61	18.0	-0.78	0.789
03	mechanical	yes	switch budget	-0.61	18.0	-0.78	0.736
04	mechanical	yes	switch budget	-1.0	18.0	-1.3	0.754
05	mechanical	yes	switch budget	-0.64	18.0	-0.89	0.755
06	mechanical	yes	switch budget	-0.335	18.0	-0.5	0.759
07	mechanical	yes	switch budget	0.445	18.0	0.56	0.88
08	mechanical	yes	switch budget	2.7	18.0	2.1	1.3
09	mechanical	yes	switch budget	5.6	18.0	5.0	1.7

hh	mode	safe	reason	bound	limit	actual	margin
10	mechanical	yes	switch budget	9.2	18.0	9.8	1.9
11	mechanical	yes	switch budget	11.8	18.0	13.5	1.7
12	mechanical	yes	switch budget	14.6	18.0	17.0	1.6
13	mechanical	yes	switch budget	17.8	18.0	19.9	1.4
14	mechanical	no	dry-bulb	20.3	18.0	22.2	1.2
15	mechanical	no	dry-bulb	22.2	18.0	23.8	1.0
16	mechanical	no	dry-bulb	22.3	18.0	22.2	1.0
17	mechanical	no	dry-bulb	22.3	18.0	21.0	1.0
18	mechanical	no	dry-bulb	22.4	18.0	20.8	1.1
19	mechanical	no	dry-bulb	20.9	18.0	20.0	1.2
20	mechanical	no	dry-bulb	19.3	18.0	18.1	1.3
21	free	yes	-	17.7	18.0	15.0	1.3
22	free	yes	-	16.9	18.0	14.1	1.1
23	free	yes	-	16.6	18.0	15.0	0.974

What would change the answer

The hours above that ran chillers on a named schedule constraint were not held by the weather. Re-planning this same day with the switch budget raised from **1 to 3** releases **14 of the 24 hours** to free cooling. Re-planning it with the minimum dwell cut from 3 h to 1 h releases **0**. On this day and this configuration the binding constraint is the number of mode changes the plant permits, not its dwell requirement and not the air.

Change to the plant's own limits	Hours released	Which hours
switch budget 1 per day to 3	+14	00:00 to 13:00
minimum dwell 3 h to 1 h	+0	none

That is a re-plan, not a promise. The switch budget exists because cycling chillers and dampers wears them out, so the hours above are available only if the plant is willing to pay that wear, and the agent deliberately does not make that trade on the operator's behalf. What it does is price it: this is the number an engineer needs in order to decide whether the setting is worth revisiting. Each of the 14 is re-planned and re-checked in `explain.py` before it may be printed here.

Generated by `AGENTIC-ARBITER/src/site_report.py` from this site's own artefacts. Every figure here is read from the file that produced it; the rows labelled `MODELED` are the only ones resting on an unmeasured assumption.